

# NLP3 PROJECT

LAB01

January 2023



nelson.vicel-farah  
antoine.zellmeyer  
karen.kaspar  
romain1.brand  
maxence.plantard

EPITA  
SCIA - Promo 2023

# Contents

<b>1</b>	<b>Language Detection (24 points) – Guided coding</b>	<b>1</b>
1.1	Question 1 (1 point) . . . . .	1
1.2	Question 2 (1 point) . . . . .	1
1.3	Question 3 (1 point) . . . . .	2
1.4	Question 4 (1 point) . . . . .	2
1.5	Question 5 (4 point) . . . . .	2
1.6	Question 6 (3 points) . . . . .	3
1.7	Question 7 (3 points) . . . . .	3
1.8	Question 8 (1 point) . . . . .	4
1.9	Question 9 (2 points) . . . . .	4
1.10	Question 10 (1 point) . . . . .	5
1.11	Question 11 (2 points) . . . . .	5
1.12	Question 12 (4 points) . . . . .	6
<b>2</b>	<b>Rotate two semantic spaces (23 points) – Not guided coding</b>	<b>6</b>
2.1	Question 1 (1 point) . . . . .	6
2.2	Question 2 (2 points) . . . . .	6
2.3	Question 3 (2 points) . . . . .	7
2.4	Question 4 (5 points) . . . . .	8
2.5	Question 5 (2 points) . . . . .	8
2.6	Question 6 (2 points) . . . . .	8
2.7	Question 7 (3 points) . . . . .	8
2.8	Question 8 (4 points) . . . . .	8
2.9	Question 9 (2 points) . . . . .	8
2.10	Question 10 (5 points) . . . . .	9
<b>3</b>	<b>Attention Exploration (22 points)</b>	<b>9</b>
3.1	Question a (2 points) . . . . .	9
3.2	Question b (4 points) . . . . .	9
3.3	Question c (5 points) . . . . .	10
3.3.1	i (2 points) . . . . .	10
3.3.2	ii (3 points) . . . . .	10
3.4	Question d (3 points) . . . . .	11
3.4.1	i (1 points) . . . . .	11
3.4.2	ii (2 points) . . . . .	12

# 1 Language Detection (24 points) – Guided coding


## Question 0 (1 point)

Try out a translation of a French sentence in Google Translate (or Bing Translate) to English. What happens if you select the wrong source language as follows? Explain in a few sentences what is happening in the backend.

The backend of the translator will process the text as if it were written in the selected language. It will consider unrecognized word as if they were names, thus restoring them as is.

### 1.1 Question 1 (1 point)

Describe the distribution of languages and give at least two comments about the dataset. (1 point)



Languages are equally distributed in the dataset.

### 1.2 Question 2 (1 point)

Do the appropriate pre-processing to maximise the accuracy of language detection. What is your strategy?

Our preprocessing includes a few steps. The first step transforms uppercase characters into lowercase characters. We then remove punctuation. For specific languages that use logograms such as Japanese and Chinese, we have chosen to apply a specific preprocessing step that adds spaces between each character. We use the same treatment for the Thai alphabet.

### 1.3 Question 3 (1 point)

**What would be the problem if your dataset was unbalanced? (1 point)**

The model would be highly biased toward languages that outnumber the other languages.

### 1.4 Question 4 (1 point)

**What techniques could you use to solve that?**

- **Oversampling:** This involves increasing the number of samples in the minority class by duplicating existing samples or generating new synthetic samples.
- **Undersampling:** This involves reducing the number of samples in the majority class by removing some of the samples.
- **Weighted loss function:** We can use a weighted loss function to give more importance to the samples in the minority class during training.
- **Stratified sampling:** This involves sampling the dataset in a way that ensures that each class is represented in the sample in proportion to its frequency in the dataset.

### 1.5 Question 5 (4 point)

**Train a model of your choice and describe the accuracy across languages. Use an 80%, 20% train-test split. Performance is not key but explain thoroughly the process and the metric(s) you are tracking.**

We chose the Multinomial Naive Bayes algorithm to classify the texts. We get an accuracy of 98%.

	precision	recall	f1-score	support
Arabic	0.99	0.99	0.99	202
Chinese	0.99	0.98	0.99	204
Dutch	0.98	0.99	0.98	198
English	1.00	0.78	0.88	257
Estonian	0.96	1.00	0.98	192
French	0.99	0.96	0.98	208
Hindi	0.96	1.00	0.98	193
Indonesian	0.97	1.00	0.98	194
Japanese	0.99	1.00	0.99	198
Korean	0.97	1.00	0.98	194
Latin	0.93	0.99	0.96	188
Persian	0.99	1.00	0.99	198
Portugese	0.96	0.99	0.98	195
Pushto	0.96	0.99	0.98	193
Romanian	0.98	0.99	0.99	197
Russian	0.99	0.99	0.99	200
Spanish	0.99	0.99	0.99	199
Swedish	1.00	1.00	1.00	201
Tamil	0.98	1.00	0.99	196
Thai	0.99	1.00	1.00	199
Turkish	0.99	1.00	0.99	198
Urdu	0.98	1.00	0.99	196
accuracy			0.98	4400
macro avg	0.98	0.98	0.98	4400
weighted avg	0.98	0.98	0.98	4400

## 1.6 Question 6 (3 points)

Train a fasttext model on Tatoeba parallel corpus and check that performance is good.

The accuracy of fasttext is 0.958 on Tatoeba. The performance is therefore good.

## 1.7 Question 7 (3 points)

Test your fasttext model on the same dataset as in question 1-5. Compare with your custom model (make sure you use the exact same data for testing). How can you explain the difference in performance between the two models?

The accuracy of fasttext on the previous dataset is way lower than our Naive Bayes model. In this case, the accuracy of fasttext is **0.80** while the Naive Bayes accuracy is **0.98**.

This difference in performance can be explained by multiple factors :

- Both model have not been trained on the same distribution. Our custom model have been trained on the same dataset it has been tested on, while fasttext has only been trained on Tatoeba.
- FastText classifiers can struggle with very high-dimensional data because they rely on dense word vectors, which can become very large as the vocabulary size increases. Naive Bayes classifiers, on the other hand, can often handle high-dimensional data well because they make independence assumptions between features.
- FastText classifiers can also struggle with very sparse data, as the dense word vectors may not be able to capture the relationships between words in the data effectively. Naive Bayes classifiers, on the other hand, can often handle sparse data well because they make independence assumptions between features.

## 1.8 Question 8 (1 point)

Compute your performance metrics yourself and compare with sklearn.

The performance metric we rely on to compare models for this kind of task is the **accuracy** which is computed by.

$$\frac{\text{Right answers}}{\text{Wrong answers} + \text{Right answers}}$$

We eventually get the same results of the `accuracy_score` function from scikit-learn.

```
# compare with the sklearn metrics
print("Accuracy: ", accuracy(y_pred, y))

# sklearn
print("Accuracy: ", accuracy_score(y_pred, y))
#####
Accuracy:  0.8009090909090909
Accuracy:  0.8009090909090909
```

## 1.9 Question 9 (2 points)

How could you improve the fasttext model performance from the previous question? Explain in a few sentences.

There are different ways to improve the fasttext model performance from the previous question:

- One way to improve the performance would be to train the fasttext model on the data it is actual tested on.
- We could also perform hyperparameter tuning. i.e we try different parameters and keep the ones that get the best performance.
- Applying preprocessing on the Tatoeba dataset.

### 1.10 Question 10 (1 point)

**Which method would you use for language detection and why?**

The most efficient method we tried seems to be the Multinomial Naive Bayes. While it depends on the context, MNB doesn't infer its responses from the assumption that words in sentences might be dependant. This means that the model focuses on more important language-specific cues : the alphabet and the lexicon. Another advantage of Naive Bayes is that it is way faster and simpler to train than fasttext.

### 1.11 Question 11 (2 points)

**Given a sentence with  $N_1$  tokens in English and  $N_2$  token in French, what would be your strategy to assign a language to such sentence?**

To determine the language of a sentence with  $N_1$  tokens in English and  $N_2$  tokens in French, we can follow the following steps:

- Identify all the words that are enclosed in quotation marks as they might indicate expressions or verbatim quotes.
- Exclude the quoted words from the count of  $N_1$  and  $N_2$
- We determine the language of the sentence by identifying the language with the most remaining tokens after the exclusion of the words between quotation. Therefore if  $N_1$  has more tokens than  $N_2$ , we consider that the sentence is written in English. Alternatively, if  $N_2$  has more tokens than  $N_1$ , we assume that the sentence is written in French.

### 1.12 Question 12 (4 points)

**Would a multilingual architecture be robust to multiple languages in a single sentence? Elaborate your answer accordingly.**

A multilingual model may be able to effectively handle multiple languages within a single sentence, depending on how it was trained and how it processes input.

There are several methods the model could use, such as identifying and processing each language separately or using a shared representation to process the entire sentence.

The model's ability to handle multiple languages in a single sentence will depend on its language identification and processing abilities, the complexity and structure of the sentence, and the diversity and representativeness of the data used to train the model.

## 2 Rotate two semantic spaces (23 points) – Not guided coding

### 2.1 Question 1 (1 point)

**Explain in a few sentences how MUSE is doing the alignment of the semantic spaces in the supervised way.**

MUSE is a tool that allows for the creation of a mapping from one language's word embeddings to another. In a supervised setting, MUSE uses a pre-existing dictionary that translates words between the two languages and the embeddings of both languages to learn the mapping. MUSE uses the Orthogonal Procrustes problem to find the optimal mapping, which involves finding an orthogonal matrix that can most accurately transform one set of embeddings into another. To do this, MUSE computes the matrix  $M$  by multiplying the matrices of the two sets of embeddings, and then uses Singular Value Decomposition to find the closest orthogonal matrix to  $M$ , which is equivalent to the desired mapping matrix.

### 2.2 Question 2 (2 points)

**What is the limit of doing that alignment based on the approach taken in the supervised way?**

There are several limitations to the supervised approach to aligning semantic spaces using MUSE:



- The approach requires a large amount of labeled parallel data, which consists of texts in both languages that are translations of one another. This can be difficult to obtain, especially for low-resource languages or specialized domains.
- The approach is limited by the quality of the parallel data. If the parallel data is noisy or contains errors, it can negatively impact the quality of the learned transformation matrix and the alignment of the semantic spaces.
- The approach is sensitive to the choice of supervised learning algorithm and hyperparameters. Different algorithms and hyperparameter choices can lead to different results, and it may be difficult to determine the optimal settings without a large amount of trial and error.
- Finally, the approach is limited to aligning the semantic spaces of two languages. If more than two languages are involved, additional transformation matrices would need to be learned and combined, which can be a complex and time-consuming process.

### 2.3 Question 3 (2 points)

**How can we align two semantic spaces in a domain specific field, e.g., in a tech company?**

To align the semantic spaces of two languages in a domain-specific field such as in a tech company, we can follow the following steps:

- Collect a large amount of parallel data that is relevant to the specific domain (in this case the tech company). The data must consist of texts in both languages that are translations of one another.
- Preprocess the data to remove noise and errors, such as typos and non-standard language.
- Create word embeddings for each language using an unsupervised approach, such as Skip-Gram or GloVe.
- Use a supervised learning algorithm, such as Procrustes analysis or Orthogonal Procrustes, to learn a transformation matrix that maps the embeddings of one language to the other using the parallel data.
- Use the transformation matrix to align the semantic spaces of the two languages by applying it to the word embeddings of one language.
- Test the alignment by comparing the meanings of words in the two languages in the common semantic space and adjusting the transformation matrix as needed.

It may also be helpful to use domain-specific data or techniques to fine-tune the alignment for the tech domain. This could include using additional parallel data that is specific to the tech domain or incorporating domain-specific knowledge into the transformation matrix.

#### 2.4 Question 4 (5 points)

Align the French space and the English space together, with the method of your choice.

#### 2.5 Question 5 (2 points)

Visualize the output on a few words of your choice. Comment on the performance of the alignment based on the output.

#### 2.6 Question 6 (2 points)

How can you find the translation of a word with this approach? Explain your method and the distance metric you choose in a few sentences.

#### 2.7 Question 7 (3 points)

Apply your approach and comment on the performance of the translation.

#### 2.8 Question 8 (4 points)

What is the limit of aligning two spaces at a sentence level? What do you suggest to improve the alignment, at a sentence level?

#### 2.9 Question 9 (2 points)

Someone, in your company, asked you to do sentiment analysis on their dataset. Given a set of sentences  $s_1, \dots, s_N$ , where  $s_i$  can be written in any language, what architecture would you use to have a vector representation of  $s_i$ ? Motivate in 2-3 bullet points.

## 2.10 Question 10 (5 points)

How would you do sentiment analysis across multiple languages in a domain specific context? Justify your approach step by step.

## 3 Attention Exploration (22 points)

### 3.1 Question a (2 points)

Describe (in one sentence) what properties of the inputs to the attention operation would result in the output  $c$  being approximately equal to  $v_j$  for some  $j \in \{1, \dots, n\}$ . Specifically, what must be true about the query  $q$ , the values  $\{v_1, \dots, v_n\}$  and/or the keys  $\{k_1, \dots, k_n\}$ ?

$$\text{with } c = \sum_{i=1}^n a_i v_i$$

$$\text{and } a_i = \frac{\exp(k_i^T q)}{\sum_{j=1}^n \exp(k_j^T q)}$$

$c$  tends to  $v_j$  when  $a_j$  tends to 1 while every other  $a_i$  tends to 0, which happens when  $k_j^T q$  is substantially higher than every other  $k_i^T q$ .

### 3.2 Question b (4 points)

Give an expression for a query vector  $q$  such that the output  $c$  is approximately equal to the average of  $v_a$  and  $v_b$ , that is,  $\frac{1}{2}(v_a + v_b)$ .

To achieve this, we need  $a_a$  and  $a_b$  to tend to  $\frac{1}{2}$  while every other  $a_i$  tend to 0. This occurs when

$$q = \lim_{C \rightarrow \infty} C * (k_a + k_b)$$

In a numerical setting,  $C$  can be an arbitrary large scalar. In this case,  $q$  is neither perpendicular to  $k_a$  nor  $k_b$  but stays perpendicular to every other vector  $k_i$ . Which gives us :

$$\begin{aligned}
a_a &= \frac{\exp(C * k_a^T(k_a + k_b))}{\exp(C * k_a^T(k_a + k_b)) + \exp(C * k_b^T(k_a + k_b)) + \sum_{i \notin \{a,b\}} \exp(C * k_i^T(k_a + k_b))} \\
&= \frac{\exp(C * k_a^T(k_a + k_b))}{\exp(C * k_a^T(k_a + k_b)) + \exp(C * k_b^T(k_a + k_b)) + (n-2)} \\
&= \frac{\exp(C * (k_a^T k_a + k_a^T k_b))}{\exp(C * (k_a^T k_a + k_a^T k_b)) + \exp(C * (k_b^T k_a + k_b^T k_b)) + (n-2)} \\
&= \frac{\exp(C * (||k_a||^2 + 0))}{\exp(C * (||k_a||^2 + 0)) + \exp(0 + C * (||k_b||^2)) + (n-2)} \\
&= \frac{\exp(C)}{2\exp(C) + n-2} \text{ (the norm of key vectors is 1)}
\end{aligned}$$

The same goes for  $a_b$ . Every other attention weight  $a_i$  can be described as follow :

$$a_i = \frac{\exp(0)}{2\exp(C) + n-2}$$

which means that  $C$  has to largely outweigh  $\log(n-2)$  in order to make  $n-2$  insignificant and allow  $a_a$  and  $a_b$  to approximate  $\frac{1}{2}$ .

### 3.3 Question c (5 points)

#### 3.3.1 i (2 points)

**Design a query  $q$  in terms of the  $\mu_i$  such that as before,  $c \approx \frac{1}{2}(v_a + v_b)$ , and provide a brief argument as to why it works.**

We can adapt the query from the previous question for this case, which gives the following query :

$$q = \lim_{C \rightarrow \infty} C * (\mu_a + \mu_b)$$

This happens due to  $a$  being vanishingly small in the covariance matrix  $\Sigma_i = aI$  which suggests that  $k_i \approx \mu_i$ .

#### 3.3.2 ii (3 points)

**When you sample  $\{k_1, \dots, k_n\}$  multiple times, and use the  $q$  vector that you defined in part i., what qualitatively do you expect the vector  $c$  will look**

like for different samples?

For a vanishingly small  $a$  we can assume that  $k_a \sim \mathcal{N}(1, \frac{1}{2})\mu_a$ . If we use the results from **question b** for a large  $C$  that makes  $(n - 2)$  insignificant we get:

$$\begin{aligned} \text{Let } X \text{ be defined by } X &\sim \mathcal{N}(1, \frac{1}{2}) \\ c &\sim \frac{\exp(C * X)}{\exp(C * X) + \exp(C)} v_a + \frac{\exp(C)}{\exp(C * X) + \exp(C)} v_b \\ &\sim \frac{1}{\exp((1 - X) * C) + 1} v_a + \frac{1}{\exp((X - 1) * C) + 1} v_b \end{aligned}$$

We can conclude that  $c$  fluctuate between  $v_a$  and  $v_b$  as we sample the key vectors multiple times.

### 3.4 Question d (3 points)

#### 3.4.1 i (1 points)

Design  $q_1$  and  $q_2$  such that  $c$  is approximately equal to  $\frac{1}{2}(v_a + v_b)$

We assume that the covariance matrices are  $\sum_i = \alpha * I$ .  
We know that the final output of the multi-headed attention is the average of each:  $c = \frac{1}{2}(c_1 + c_2)$ , therefore we need to have  $c_1 \approx v_a$  and  $c_2 \approx v_b$ .  
Let  $C$  be an arbitrary large scalar. Let  $q_1 = C * \mu_a$  and  $q_2 = C * \mu_b$ .  
We hence have:

$$\begin{aligned} c_1 &= \sum_{i=1}^n \alpha_i * v_i \\ \text{if } i \text{ is different from } a \text{ we have:} \\ \alpha_i &\approx \frac{\exp(0)}{\exp(C)} \approx \frac{1}{\exp(C)} \approx 0 \\ \text{otherwise :} \\ \alpha_a &\approx \frac{\exp(C)}{\exp(C)} \approx 1 \\ c_1 &= \sum_{i=1}^n \alpha_i * v_i \end{aligned}$$

Therefore all the elements of the sum are cancelled except for the element at the index  $a$ : hence  $c_1 \approx v_a$

In the same manner  $c_2 \approx v_b$

So  $c \approx \frac{1}{2}(v_a + v_b)$

### 3.4.2 ii (2 points)

Take the query vectors  $q_1$  and  $q_2$  that you designed in part i. What, qualitatively, do you expect the output  $c$  to look like across different samples of the key vectors? Please briefly explain why. You can ignore cases in which  $q_i^T k_a < 0$ .

For a vanishingly small  $a$  we can write that  $k_a \sim \mathcal{N}(1, \frac{1}{2})\mu_a$  while for any other  $i$ ,  $k_i = \mu_i$ .

Let  $X$  be defined by  $X \sim \mathcal{N}(1, \frac{1}{2})$   
The query vectors that we designed before are:

$$q_1 = C * \mu_a \text{ and } q_2 = C * \mu_b$$

$$c_1 = \sum_{i=1}^n \alpha_i * v_i$$

We have seen that the value is cancelled for all indexes except when  $i = a$ .

For  $i = a$ , we have:

$$\begin{aligned} c_1 &\sim \frac{\exp(k_a^T q_1)}{\exp(k_a^T q_1)} v_a \\ c_1 &\sim \frac{\exp(X * \mu_a^T * C * \mu_a)}{\exp(X * \mu_a^T * C * \mu_a)} \approx \frac{\exp(X * C)}{\exp(X * C)} v_a \\ &\approx v_a \end{aligned}$$

Therefore  $c_1 \approx v_a$  and in the same manner  $c_2 \approx v_b$

For  $i = b$ , and  $c_2$  we have

$$\begin{aligned} c_2 &\sim \frac{\exp(k_b^T q_2)}{\exp(k_b^T q_2) + \exp(k_a^T q_2)} v_b + \frac{\exp(k_a^T q_2)}{\exp(k_b^T q_2) + \exp(k_a^T q_2)} v_a \\ c_2 &\sim \frac{\exp(\mu_b^T * C * \mu_b)}{\exp(\mu_b^T * C * \mu_b)} \approx \frac{\exp(C)}{\exp(C)} v_b \approx v_b \end{aligned}$$

We end up with the same result as in the first question with  $c \approx \frac{1}{2}(v_a + v_b)$  which means that  $c$  will approach this value as we sample the key vectors